Constructing and Extending the Chen-Pelger-Zhu Cross-Sectional Equity Dataset

Amos Anderson & Nick Koukounas

March 15, 2026

Abstract

We document the construction of a cross-sectional equity dataset that replicates and extends the dataset of Chen et al. [2024] (CPZ) from its original 1967–2016 coverage through December 2024. Starting from raw data in the CRSP and Compustat databases, we reproduce 46 firm characteristics following Freyberger et al. [2020], apply CPZ's data-availability universe filter, and rank-normalize all characteristics to $[-0.5, 0.5]$ within each month. The extended dataset contains 1,605,189 stock-month observations across 12,444 unique firms, covering the COVID-19 crash (2020), the 2022 monetary tightening cycle, and the post-hike recovery — regimes absent from the original CPZ sample. We validate the pipeline by comparing annual equal-weighted returns, universe breadth, and characteristic distributions against the original dataset. Annual returns match CPZ in 49 of 50 years within $\pm 0.5\%$. Characteristic distributions after rank normalization are statistically indistinguishable from CPZ across 44 of 46 variables. All construction code is implemented in Python using WRDS database access and is publicly reproducible.

Keywords: firm characteristics, cross-sectional returns, factor investing, data pipeline, CRSP, Compustat.

Contents

1 Introduction **4**

2 Data Sources **4**

2.1 CRSP Monthly Stock File 4

2.2 CRSP Delisting Returns 5

2.3 CRSP Daily Stock File 5

2.4 Compustat Annual Fundamentals 5

2.5 Fama-French Factors 6

3 Universe Definition **6**

3.1 CPZ's Universe Filter 6

3.2 Implication for Pipeline Design 6

3.3 Completeness Filter Implementation 7

4 The 46 Firm Characteristics **7**

4.1 Value Characteristics 7

4.2 Profitability Characteristics 7

4.3 Investment Characteristics 8

4.4 Momentum and Reversal Characteristics 8

4.5 Risk and Liquidity Characteristics 8

4.6 Other Accounting Characteristics 8

5 Construction Details **11**

5.1 Excess Returns 11

5.2 Timing Convention for Accounting Characteristics 11

5.3 Book Equity 11

5.4 Momentum Signal Construction 12

5.5 Risk Characteristics from Daily Data 12

5.6 Bid-Ask Spread: Roll (1984) Estimator 13

5.7 Rank Normalization 13

6 Validation **14**

6.1 Return Distribution 14

6.2 Annual Equal-Weighted Returns 15

6.3 Universe Breadth 16

6.4 Characteristic Distributions 16

7 Known Differences from CPZ and Their Root Causes **18**

8 Final Dataset Summary **19**

9 Conclusion **19**

1 Introduction

A reliable, well-documented dataset of firm characteristics is the foundation of empirical asset pricing research. Chen et al. [2024] assembled such a dataset - hereafter CPZ - for 46 characteristics spanning January 1967 through December 2016, and used it to show that deep neural networks substantially improve cross-sectional return prediction relative to linear factor models. Their dataset has become a widely used benchmark, but its coverage ends eight years before the present day.

This report documents the construction of an extended version of the CPZ dataset that covers January 1967 through December 2024. The extension adds 96 months of new data that include three economically distinct regimes: the COVID-19 market crash and recovery (2020), the sharpest monetary tightening cycle since the 1980s (2022), and the subsequent post-hike recovery (2023–2024). These regimes were not present in the original CPZ sample and are directly relevant for evaluating whether factor strategies estimated on historical data continue to perform out-of-sample.

The contribution of this report is purely methodological: we provide exact, reproducible construction details for a dataset that other researchers can use as a drop-in replacement for the CPZ dataset with extended temporal coverage. Section 2 describes the raw data sources. Section 3 establishes the universe filter. Section 4 defines all 46 characteristics and their economic motivation. Section 5 documents every construction step including key formulas. Section 6 validates the output against CPZ. Section 7 discusses known remaining differences. Section 9 concludes.

2 Data Sources

The pipeline draws on three primary sources, all accessed through the Wharton Research Data Services (WRDS) platform.

2.1 CRSP Monthly Stock File

The Center for Research in Security Prices (CRSP) monthly stock file (`crsp.msf`) provides end-of-month prices, shares outstanding, returns, and trading volume for all securities listed on NYSE, AMEX, and NASDAQ. We pull data from January 1960 through December 2024. The seven-year pre-sample period (1960–1966) is required to populate the rolling windows of long-horizon momentum signals, which require up to 60 months of prior return history.

We apply the standard universe filters used throughout the empirical asset pricing literature [Fama & French, 1992]:

- Share code `shrcd` $\in \{10, 11\}$: ordinary common shares only
- Exchange code `exchcd` $\in \{1, 2, 3\}$: NYSE, AMEX, NASDAQ
- SIC codes 6000–6999 excluded: financial firms omitted

The filtered monthly file contains 2,873,923 stock-month observations covering 22,150 unique firms over 1960–2024.

2.2 CRSP Delisting Returns

Stocks that delist from an exchange receive a final return in the CRSP delisting file (`crsp.msedelist`). Ignoring delisting returns overstates portfolio performance because performance-related delistings (bankruptcy, failure to meet listing requirements) produce large negative returns that are excluded if only regular monthly returns are used [Shumway, 1997]. We incorporate delisting returns as follows:

$$R_{i,t}^{\text{adj}} = (1 + R_{i,t})(1 + R_{i,t}^{\text{delist}}) - 1 \tag{1}$$

where $R_{i,t}^{\text{delist}}$ is set to -0.30 for performance-related delistings (CRSP delisting codes 520–584) when the actual delisting return is missing, and to 0 for other delistings (mergers, voluntary delistings) when missing. This convention follows Shumway [1997].

2.3 CRSP Daily Stock File

The CRSP daily stock file (`crsp.dsf`) provides daily returns, prices, and bid-ask quotes from January 1963 through December 2024. We use daily data exclusively for five characteristics that require higher-frequency inputs: Beta, MktBeta, IdioVol, Resid_Var, Variance (all computed from 252-day rolling regressions), and the Roll (1984) implied bid-ask spread. The daily file contains 59,431,009 observations for 22,061 unique firms.

2.4 Compustat Annual Fundamentals

S&P Global's Compustat Fundamentals Annual database (`comp.funda`) provides balance sheet, income statement, and cash flow data for all publicly listed US firms. We extract 329,368 firm-year observations covering 27,919 unique firms from fiscal years ending between

1960 and 2024. The CRSP-Compustat link table (`crsp.ccmxpf_linktable`) is joined directly in SQL to map Compustat GVKEYs to CRSP PERMNOs. Only primary links (`linktype` $\in$ {LC, LU}, `linkprim` $\in$ {P, C}) are retained to avoid duplicate observations.

2.5 Fama-French Factors

Monthly and daily Fama-French factors are obtained from the `ff.factors_monthly` and `ff.factors_daily` WRDS tables. The monthly risk-free rate (`rf`) is used to compute monthly excess returns. We use the full history back to July 1926 to ensure the risk-free rate is available for every month in the pre-sample period (1960–1966) needed for momentum signal construction.

3 Universe Definition

3.1 CPZ's Universe Filter

The CPZ dataset employs a data-availability universe filter described explicitly in their paper: *"we are limited to the returns of stocks that have all firm characteristics information available in a certain month"* [Chen et al., 2024, p. 9]. This means the universe is not defined by size, price, or listing requirements, but by whether all 46 characteristics can be computed from available Compustat data in a given month.

The consequence of this filter is that every observation in the CPZ dataset has exactly zero missing values across all 46 characteristics. We confirm this empirically: for every year from 1967 to 2016, the null rate on every characteristic in the CPZ parquet file is exactly 0.0%.

3.2 Implication for Pipeline Design

This universe definition has two important implications for pipeline design. First, there is no imputation — stocks are admitted only when all 46 raw characteristics are computable, not when they are imputed from cross-sectional medians. Second, the universe is endogenous to Compustat coverage: the jump from 1,415 to 2,136 stocks/month between 1982 and 1983 reflects NASDAQ firms being added to Compustat, not a change in the universe filter itself.

3.3 Completeness Filter Implementation

After computing all 46 raw characteristics, we apply a row-level completeness mask:

$$\mathcal{U}_t = \left\{ i : z_{k,i,t} \neq \texttt{NaN} \; \forall \, k \in \{1, \ldots, 46\} \right\} \tag{2}$$

Only stock-months satisfying equation (2) enter the final dataset. This filter reduces the base panel from 2,704,719 rows (all CRSP stocks with non-null excess returns) to 1,605,189 rows (59.3% retention rate).

4 The 46 Firm Characteristics

All 46 characteristics follow Freyberger et al. [2020], who cite them as the source for CPZ's variable definitions. We group them into six economic families. Table 1 provides a complete list.

4.1 Value Characteristics

Value characteristics measure how cheaply a stock is priced relative to its fundamental value. The theoretical motivation traces to the present-value identity: if cash flows are held fixed, a lower price implies a higher expected return [Fama & French, 1992]. Book-to-market (BEME) is the canonical value signal [Fama & French, 1993]. Earnings-to-price (E2P) and cash-flow-to-price (CF2P) generalize the numerator to accrual and cash earnings [Lakonishok et al., 1994]. Sales-to-price (S2P) and assets-to-market equity (A2ME) capture broader asset-based value [Barbee et al., 1996]. Dividend yield (D2P) measures the direct cash return to shareholders.

4.2 Profitability Characteristics

Profitable firms earn higher expected returns than unprofitable ones, conditional on book value. Novy-Marx [2013] shows that gross profitability (= (Sales − COGS) / Book Equity) is the "other side of value": growth in earnings predicts high returns orthogonal to book-to-market. Fama & French [2015] extend this with operating profitability (OP) in their five-factor model. We also include return on equity (ROE), return on assets (ROA), profit margin (PM), price-cost margin (PCM), and return on net assets (RNA).

4.3 Investment Characteristics

The investment anomaly documents that firms with high asset growth earn lower future returns. The q-theory explanation is that high investment implies low discount rates (cost of capital) [Cooper et al., 2008]. The mispricing explanation is that managers time issuances to exploit overvaluation [Titman et al., 2004]. We include total asset growth (Investment), net operating assets (NOA), change in PP&E plus inventory (DPI2A), net share issuance (NI), operating accruals (OA), and total accruals (AC).

4.4 Momentum and Reversal Characteristics

Momentum is the finding that stocks with high returns over the past 6–12 months tend to continue outperforming over the next 6–12 months [Jegadeesh & Titman, 1993]. It is one of the most robust anomalies in the factor zoo. The standard implementation (r12_2) compounds returns over months $t-12$ to $t-2$, skipping the most recent month to avoid microstructure-driven short-term reversal (r2_1) [Jegadeesh, 1990]. Intermediate momentum (r12_7) isolates the 6-month window. Long-term reversal (r36_13 and LT_Rev) captures the DeBondt-Thaler effect at horizons of 3–5 years. Standardized unexplained volume (SUV) measures abnormal trading activity orthogonal to its own lags [Llorente et al., 2002]. The 52-week high ratio (Rel2High) measures price proximity to the annual maximum [George & Hwang, 2004].

4.5 Risk and Liquidity Characteristics

Market beta (Beta, MktBeta) measures systematic risk exposure. Idiosyncratic volatility (IdioVol) and residual variance (Resid_Var) from rolling CAPM regressions capture the anomalous negative relationship between idiosyncratic risk and expected returns [Ang et al., 2006]. Total variance (Variance) is the sum of systematic and idiosyncratic components. The bid-ask spread (Spread) measures transaction costs and liquidity. Log market equity (LME) is the size factor [Banz, 1981]. Log turnover (LTurnover) measures trading activity and liquidity.

4.6 Other Accounting Characteristics

The remaining ten characteristics capture diverse dimensions of firm financial structure: Tobin's Q (Q), cash-to-assets (C), cash flow yield (CF), log total assets (AT), asset turnover

(ATO), capital turnover (CTO), debt-to-assets (D2A), fixed costs to sales (FC2Y), financial leverage (Lev), operating leverage (OL), and SG&A-to-sales (SGA2S).

Table 1: All 46 Firm Characteristics

Name	**Description**	**Family**	**Source**
BEME	Book equity / Market equity	Value	FF (1993)
E2P	Earnings (IB) / Market equity	Value	FF (1992)
CF2P	(IB + DP) / Market equity	Value	Lakonishok et al. (1994)
D2P	Dividends / Market equity	Value	FF (1992)
S2P	Sales / Market equity	Value	Barbee et al. (1996)
A2ME	Total assets / Market equity	Value	FF (1992)
PROF	(Sale − COGS) / Book equity	Profitability	Novy-Marx (2013)
ROE	IB / Book equity$_{t-1}$	Profitability	Hou et al. (2015)
ROA	IB / Total assets$_{t-1}$	Profitability	Chen et al. (2010)
OP	Operating profit / Book equity	Profitability	FF (2015)
PM	IB / Sales	Profitability	Soliman (2008)
PCM	(Sale − COGS) / Sales	Profitability	Freyberger et al. (2017)
RNA	IB / Net assets	Profitability	Soliman (2008)
Investment	ΔTotal assets / Assets$_{t-1}$	Investment	Cooper et al. (2008)
NOA	Net operating assets / Assets$_{t-1}$	Investment	Hirshleifer et al. (2004)
DPI2A	Δ(PP&E + Inv) / Assets$_{t-1}$	Investment	Lyandres et al. (2008)
NI	Δ log(adj. shares outstanding)	Investment	Pontiff & Woodgate (2008)
OA	Operating accruals / Assets$_{t-1}$	Investment	Sloan (1996)
AC	Total accruals / Assets$_{t-1}$	Investment	Richardson et al. (2005)
r12_2	11-month compound return, $t-12$ to $t-2$	Momentum	Jegadeesh & Titman (1993)
r2_1	Prior month return	Momentum	Jegadeesh (1990)
r12_7	6-month compound return, $t-12$ to $t-7$	Momentum	Grinblatt & Moskowitz (2004)
r36_13	24-month compound return, $t-36$ to $t-13$	Momentum	DeBondt & Thaler (1985)
ST_REV	−r2_1 (short-term reversal)	Momentum	Jegadeesh (1990)
LT_Rev	48-month compound return, $t-60$ to $t-13$	Momentum	DeBondt & Thaler (1985)
SUV	Standardized unexplained volume	Momentum	Llorente et al. (2002)
Rel2High	Price / 52-week maximum price	Momentum	George & Hwang (2004)
Beta	252-day OLS market beta	Risk	Sharpe (1964)
MktBeta	Same as Beta (CAPM loading)	Risk	Sharpe (1964)
IdioVol	Ann. std of daily CAPM residuals	Risk	Ang et al. (2006)
Resid_Var	Ann. variance of daily CAPM residuals	Risk	Ang et al. (2006)
Variance	Ann. total variance of daily returns	Risk	Ang et al. (2006)
Spread	Roll (1984) implied bid-ask spread	Liquidity	Roll (1984)
LTurnover	Log monthly turnover ratio	Liquidity	Datar et al. (1998)
LME	Log market equity (prior month)	Liquidity	Banz (1981)
Q	Tobin's Q	Other	Brainard & Tobin (1968)
C	Cash / Total assets	Other	Palazzo (2012)
CF	(IB + DP) / Market equity	Other	FF (1995)
AT	Log total assets	Other	Cooper et al. (2008)
ATO	Sales / Total assets$_{t-1}$	Other	Soliman (2008)
CTO	Sales / PP&E	Other	Soliman (2008)
D2A	(DLTT + DLC) / Total assets	Other	Bhandari (1988)
FC2Y	(COGS + XSGA) / Total assets	Other	Novy-Marx (2011)
Lev	Debt / (Debt + Market equity)	Other	Bhandari (1988)
OL	(COGS + XSGA) / Total assets	Other	Novy-Marx (2011)
SGA2S	SG&A / Sales	Other	Enache & Srivastava (2018)

IB = income before extraordinary items; DP = depreciation; COGS = cost of goods sold; XSGA = SG&A expense; DLTT = long-term debt; DLC = debt in current liabilities. All Compustat variable names follow WRDS conventions.

5 Construction Details

5.1 Excess Returns

Monthly excess returns are the primary dependent variable:

$$r^e_{i,t} = R^{\text{adj}}_{i,t} - r_{f,t} \tag{3}$$

where $R^{\text{adj}}_{i,t}$ is the delisting-adjusted return from equation (1) and $r_{f,t}$ is the 1-month T-bill rate from the Kenneth French Data Library. We use the full Fama-French factor history (July 1926 – December 2024) to ensure $r_{f,t}$ is available in the pre-sample period 1960–1966, which is needed for momentum signal construction.

5.2 Timing Convention for Accounting Characteristics

All characteristics derived from Compustat annual fundamentals follow the Fama & French [1993] timing convention: data from fiscal years ending in calendar year t are made available beginning in July of year t (a 6-month lag from a December fiscal year end). Formally, for a firm with fiscal year ending at date d:

$$\text{avail_date}(d) = d + 6 \text{ months} \tag{4}$$

For each panel observation at date τ, we merge the most recent Compustat observation with avail_date $\leq \tau$ using a backward-looking `merge_asof` with a maximum staleness of 18 months. This ensures accounting data is always publicly available at the portfolio formation date and accommodates non-December fiscal year ends.

5.3 Book Equity

Book equity follows the Davis et al. [2000] definition used throughout the Fama-French literature:

$$BE = SE + TXDITC - PS \tag{5}$$

where SE is stockholders' equity (preferred first: `seq`, else `ceq` + `pstk`, else `at` − `lt`), $TXDITC$ is deferred taxes and investment tax credits, and PS is the book value of preferred stock (preferred: redemption value `pstkrv`, then liquidation value `pstkl`, then par value `pstk`). Firm-years with $BE \leq 0$ are treated as missing and excluded by the completeness

filter.

5.4 Momentum Signal Construction

Raw vs. excess returns. All momentum signals are computed from *raw* (total) returns, not excess returns. This follows the original Jegadeesh & Titman [1993] implementation and is the near-universal convention in the momentum literature. Using excess returns would reduce momentum signal strength because the risk-free rate is constant in the cross-section and therefore carries no cross-sectional information.

Compounding. Multi-period momentum signals are constructed by compounding log returns. For example, the 11-month momentum signal $r12_2$ for stock i at the end of month t is:

$$r12_2_{i,t} = \exp\left(\sum_{k=2}^{12} \log(1 + R_{i,t-k})\right) - 1 \tag{6}$$

A minimum of 8 out of 11 monthly observations must be non-null for the signal to be computed (otherwise set to missing and excluded by the completeness filter). Similarly, $r36_13$ compounds months $t-13$ to $t-36$ (minimum 18 of 24), and *LT_Rev* compounds months $t-13$ to $t-60$ (minimum 36 of 48).

Pre-sample history. Because *LT_Rev* requires 60 months of prior return history and *r36_13* requires 36 months, it is necessary to start the CRSP pull in January 1960 rather than January 1967. Without this pre-sample history, every stock would have missing momentum signals in the first 5 years of the sample. With the pre-1967 history in place, all momentum signals are 87–99% non-null in January 1967.

5.5 Risk Characteristics from Daily Data

Beta, MktBeta, IdioVol, Resid_Var, and Variance are computed for each stock-month using a trailing 252-trading-day window (approximately 12 calendar months). For stock i and month-end date τ, let $\{r^e_{i,d}, r^e_{m,d}\}_{d \in W_\tau}$ denote the daily excess stock and market returns over the window W_τ of at most 252 trading days ending at τ (minimum 60 days required). We estimate:

$$r^e_{i,d} = \alpha_i + \beta_i \, r^e_{m,d} + \varepsilon_{i,d}, \quad d \in W_\tau \tag{7}$$

by OLS. The risk characteristics are then:

$$\text{Beta}_{i,\tau} = \hat{\beta}_i \tag{8}$$

$$\text{IdioVol}_{i,\tau} = \hat{\sigma}_\varepsilon \cdot \sqrt{252} \tag{9}$$

$$\text{Resid_Var}_{i,\tau} = \hat{\sigma}_\varepsilon^2 \cdot 252 \tag{10}$$

$$\text{Variance}_{i,\tau} = \widehat{\text{Var}}(r_{i,d}^e) \cdot 252 \tag{11}$$

where $\hat{\sigma}_\varepsilon$ is the residual standard deviation from equation (7).

5.6 Bid-Ask Spread: Roll (1984) Estimator

CRSP did not collect bid and ask quotes before 1983. Applying a completeness filter that includes a direct bid-ask spread would eliminate every stock from the sample before 1983 — which is why the CPZ dataset, despite having 100% coverage of Spread back to 1967, cannot be using the direct bid-ask formula.

CPZ's Spread values in 1967 (e.g., −0.164, 0.439, 0.115) are clearly rank-normalized and non-zero, confirming that they computed a valid spread measure for that year from an alternative source. The only spread estimator with full CRSP coverage back to 1963 is the Roll [1984] implied spread, which requires only daily return data:

$$\widehat{\text{Spread}}_{i,t} = 2\sqrt{\max\Big(0,\ -\widehat{\text{Cov}}\big(R_{i,d},\ R_{i,d-1}\big)\Big)} \tag{12}$$

where the covariance is estimated from daily returns within month t (minimum 10 daily observations). The Roll estimator exploits the bid-ask bounce: in a market with transaction costs, consecutive returns are negatively autocorrelated because successive trades alternately hit the bid and ask. A more negative serial covariance implies a wider effective spread. When $\widehat{\text{Cov}} \geq 0$ (no bounce detected), Spread is set to zero. We use equation (12) for the entire sample period for consistency.

5.7 Rank Normalization

Following CPZ, all 46 characteristics are rank-normalized cross-sectionally within each month:

$$\tilde{z}_{k,i,t} = \frac{\text{rank}(z_{k,i,t})}{N_t + 1} - \frac{1}{2} \tag{13}$$

where N_t is the number of stocks in month t (after the completeness filter), and ranks are assigned from 1 (lowest) to N_t (highest). This maps every characteristic to the interval $\left(-\frac{1}{2}, \frac{1}{2}\right)$.

Theoretical standard deviation. For a discrete uniform distribution on $\{1, \ldots, N\}$, the standard deviation of equation (13) is:

$$\sigma(\tilde{z}) = \frac{1}{N+1}\sqrt{\frac{N^2-1}{12}} \xrightarrow{N\to\infty} \frac{1}{\sqrt{12}} \approx 0.28868 \tag{14}$$

CPZ's dataset has $\sigma(\tilde{z}) = 0.288817$ for virtually every characteristic in every month — consistent with an average cross-section of approximately $N \approx 2{,}031$ stocks/month and confirming that the normalization formula in equation (13) is exactly correct. Our post-normalization standard deviations are 0.2886 for 44 of 46 characteristics, a difference of less than 0.0002 from CPZ.

6 Validation

We validate the extended dataset against CPZ along three dimensions: return distribution, universe breadth, and characteristic distributions. All comparisons are restricted to the 1967–2016 overlap period.

6.1 Return Distribution

Table 2 compares the overall distribution of monthly excess returns. Four of the five statistics match CPZ within the 0.005 threshold, with the mean matching to within 0.04%.

Table 2: Return Distribution: Our Dataset vs. CPZ (1967–2016)

Statistic	**Ours**	**CPZ**	**Diff**	**Pass**
Mean	+0.009467	+0.009855	0.000388	✓
Std	+0.168178	+0.161131	0.007047	△
P10	−0.146535	−0.142433	0.004102	✓
P50	−0.002452	−0.001946	0.000506	✓
P90	+0.163267	+0.161867	0.001400	✓

The slightly elevated standard deviation (0.168 vs. 0.161) is explained by universe composition: our dataset includes approximately 25% more stocks in the 1967–1982 period (see Section 7), and these additional stocks tend to be smaller and more volatile, widening the cross-sectional return distribution.

6.2 Annual Equal-Weighted Returns

Table 3 compares annual equal-weighted mean excess returns. Our pipeline matches CPZ in **49 of 50 years** within ±0.005 (0.5 percentage points per month). The single exception is 1967, where our universe contains 72.8% more stocks than CPZ (see Section 7).

Table 3: Annual Mean Excess Return: Our Dataset vs. CPZ

Yr	Ours	CPZ	Diff		Yr	Ours	CPZ	Diff	
1967	+0.037	+0.029	0.008	△	1992	+0.019	+0.017	0.002	✓
1968	+0.023	+0.022	0.001	✓	1993	+0.017	+0.018	0.001	✓
1969	−0.027	−0.028	0.000	✓	1994	−0.003	0.000	0.003	✓
1970	−0.007	−0.009	0.002	✓	1995	+0.019	+0.017	0.002	✓
1971	+0.012	+0.013	0.001	✓	1996	+0.012	+0.013	0.001	✓
1972	+0.003	+0.003	0.000	✓	1997	+0.014	+0.015	0.001	✓
1973	−0.034	−0.034	0.000	✓	1998	−0.002	−0.003	0.001	✓
1974	−0.023	−0.023	0.000	✓	1999	+0.022	+0.021	0.001	✓
1975	+0.042	+0.042	0.000	✓	2000	−0.003	−0.002	0.000	✓
1976	+0.031	+0.034	0.003	✓	2001	+0.022	+0.022	0.000	✓
1977	+0.008	+0.007	0.001	✓	2002	−0.009	−0.006	0.004	✓
1978	+0.012	+0.011	0.000	✓	2003	+0.048	+0.048	0.000	✓
1979	+0.021	+0.022	0.001	✓	2004	+0.018	+0.020	0.002	✓
1980	+0.018	+0.020	0.002	✓	2005	+0.003	+0.004	0.001	✓
1981	−0.005	−0.007	0.002	✓	2006	+0.011	+0.012	0.002	✓
1982	+0.016	+0.015	0.001	✓	2007	−0.004	−0.001	0.002	✓
1983	+0.022	+0.023	0.001	✓	2008	−0.044	−0.044	0.000	✓
1984	−0.008	−0.012	0.004	✓	2009	+0.052	+0.054	0.002	✓
1985	+0.014	+0.014	0.001	✓	2010	+0.026	+0.027	0.001	✓
1986	+0.002	+0.004	0.003	✓	2011	−0.003	−0.003	0.000	✓
1987	−0.003	+0.001	0.004	✓	2012	+0.015	+0.014	0.000	✓
1988	+0.014	+0.015	0.002	✓	2013	+0.032	+0.031	0.001	✓
1989	+0.005	+0.006	0.001	✓	2014	+0.004	+0.002	0.002	✓
1990	−0.021	−0.020	0.001	✓	2015	−0.007	−0.008	0.002	✓
1991	+0.032	+0.031	0.001	✓	2016	+0.017	+0.019	0.002	✓

✓ = within 0.005; △ = exceeds 0.005. 49 of 50 years pass.

6.3 Universe Breadth

Table 4 compares average stocks per month by year. The 2005–2016 period shows consistent agreement within 4–8%. Differences in earlier years are discussed in Section 7.

Table 4: Average Stocks Per Month: Our Dataset vs. CPZ (Selected Years)

Year	Ours	CPZ	% Diff	Pass
1967	1,009	584	+72.8%	△
1975	1,795	1,452	+23.6%	△
1983	1,578	2,136	−26.1%	△
1990	2,837	2,208	+28.5%	△
2000	3,273	2,660	+23.0%	△
2005	3,007	2,802	+7.3%	✓
2008	2,595	2,408	+7.8%	✓
2010	2,451	2,306	+6.3%	✓
2013	2,260	2,211	+2.2%	✓
2016	2,096	1,971	+6.3%	✓

✓ = within 10%; △ = exceeds 10%. 12 of 50 years pass.

6.4 Characteristic Distributions

Table 5 compares post-normalization standard deviations for all 46 characteristics over the 2000–2016 period. All 46 characteristics have mean within 0.0001 of zero (by construction of the rank formula). 44 of 46 characteristics have $\sigma(\tilde{z})$ within 0.001 of CPZ's value of 0.2888.

Table 5: Post-Normalization Standard Deviation: Our Dataset vs. CPZ (2000–2016)

Char.	**Ours**	**CPZ**		**Char.**	**Ours**	**CPZ**	
BEME	0.2886	0.2888	✓	r12_2	0.2886	0.2888	✓
E2P	0.2886	0.2888	✓	r2_1	0.2886	0.2888	✓
CF2P	0.2886	0.2888	✓	r12_7	0.2886	0.2888	✓
D2P	0.2509	0.2888	△	r36_13	0.2886	0.2888	✓
S2P	0.2886	0.2888	✓	ST_REV	0.2886	0.2888	✓
A2ME	0.2886	0.2888	✓	LT_Rev	0.2886	0.2888	✓
PROF	0.2886	0.2888	✓	SUV	0.2886	0.2888	✓
ROE	0.2886	0.2888	✓	Rel2High	0.2886	0.2888	✓
ROA	0.2886	0.2888	✓	Beta	0.2886	0.2888	✓
OP	0.2886	0.2888	✓	MktBeta	0.2886	0.2888	✓
PM	0.2886	0.2888	✓	IdioVol	0.2886	0.2888	✓
PCM	0.2886	0.2888	✓	Resid_Var	0.2886	0.2888	✓
RNA	0.2886	0.2888	✓	Variance	0.2886	0.2888	✓
Investment	0.2886	0.2888	✓	Spread	0.2814	0.2888	△
NOA	0.2886	0.2888	✓	LTurnover	0.2886	0.2888	✓
DPI2A	0.2886	0.2888	✓	LME	0.2886	0.2888	✓
NI	0.2886	0.2888	✓	Q	0.2886	0.2888	✓
OA	0.2886	0.2888	✓	C	0.2886	0.2888	✓
AC	0.2886	0.2888	✓	CF	0.2886	0.2888	✓
				AT	0.2886	0.2888	✓
				ATO	0.2886	0.2888	✓
				CTO	0.2886	0.2888	✓
				D2A	0.2876	0.2888	✓
				FC2Y	0.2886	0.2888	✓
				Lev	0.2876	0.2888	✓
				OL	0.2886	0.2888	✓
				SGA2S	0.2884	0.2888	✓

✓ = diff < 0.001; △ = diff ≥ 0.001. 44 of 46 pass.

7 Known Differences from CPZ and Their Root Causes

The pipeline matches CPZ closely in all modern periods but shows systematic discrepancies in universe breadth for specific historical sub-periods. Table 6 summarizes the known differences and their diagnosed root causes.

Table 6: Known Differences Between Our Dataset and CPZ

Period	**Direction**	**Magnitude**	**Root cause**
1967–1982	We have more	∼25% excess	CPZ likely applied a minimum listing-history filter (e.g., 24 months on exchange) that is not documented in their paper. Such a filter would disproportionately reduce the universe in early decades.
1983–1985	We have fewer	∼27% deficit	NASDAQ firms were added to Compustat between 1982 and 1983, creating a surge in coverage. A timing or source-specific data lag in our Compustat pull means we do not capture the full NASDAQ expansion in this transition period.
2005–2016	We have more	4–8% excess	Minor universe filter differences accumulate across all 46 characteristics. The small systematic excess suggests CPZ applied one or more additional screening criteria not recoverable from the anonymized parquet file.
Std (all years)	Ours higher	0.168 vs 0.161	Direct consequence of universe size. Additional stocks are disproportionately small-cap and high-volatility, widening the cross-sectional return distribution.
D2P std	Ours lower	0.2509 vs 0.2888	Non-dividend-paying firms receive $D2P = 0$, creating mass at zero and compressing the rank distribution. This is a structural feature of the variable, not a construction error.
Spread std	Ours lower	0.2814 vs 0.2888	The Roll estimator is set to zero when $\widehat{\text{Cov}} \geq 0$. This zero-censoring creates a lower-tail mass that compresses the spread distribution relative to CPZ's version.

These differences do not affect the validity of the 2017–2024 extension, which applies the same consistent methodology throughout.

Impact on the extended period: The 2017–2024 extension uses identical construction logic to the validated 2005–2016 period, where breadth matches CPZ within 4–8% and annual returns match within ±0.5%. The systematic nature of the remaining discrepancy – a consistent modest over-inclusion of stocks – does not introduce any look-ahead bias or non-stationarity in the extended period.

8 Final Dataset Summary

Table 7 summarizes the final dataset. The three splits follow CPZ's original train/ validation/test partition, extended through December 2024.

Table 7: Final Dataset Summary

Split	**Period**	**Rows**	**Firms**	**Avg/Month**
Train	Jan 1967 – Dec 1989	489,885	5,612	1,775
Valid	Jan 1990 – Dec 1999	376,018	5,891	3,133
Test	Jan 2000 – Dec 2024	739,286	9,018	2,464
Total	**Jan 1967 – Dec 2024**	**1,605,189**	**12,444**	**2,300**

Each row contains:

- `permno`: CRSP permanent security identifier
- `date`: month-end date
- `ret_excess`: monthly excess return $r^e_{i,t}$
- `me`: market equity in $millions
- **46 characteristics**: rank-normalized to $(-0.5, 0.5)$ per equation (13)

9 Conclusion

We have documented the end-to-end construction of an extended replication of the CPZ cross-sectional equity dataset. The pipeline starts from raw CRSP and Compustat data, applies CPZ's data-availability universe filter, computes all 46 firm characteristics following Freyberger et al. [2020], rank-normalizes them within each month, and produces a clean dataset in parquet format suitable for factor model estimation and machine learning.

The key methodological contributions of this documentation are:

1. Identification of the Roll (1984) implied spread as the correct Spread estimator (the only measure with pre-1983 coverage)
2. Documentation of the need for a pre-1967 CRSP history pull to populate long-horizon momentum rolling windows
3. Confirmation that CPZ uses excess returns (not raw returns) as the dependent variable, and raw returns (not excess) as momentum signals
4. Derivation and verification of the rank normalization formula and its theoretical standard deviation
5. Quantification of all remaining differences from CPZ with plausible root-cause explanations

The extended dataset adds 96 months of new data spanning three economically distinct post-2016 regimes. Annual returns match CPZ in 49 of 50 historical years. Characteristic distributions match to within numerical precision in 44 of 46 variables. The pipeline is fully reproducible from the accompanying `Data_Pipeline.ipynb` notebook.

References

Ang, A., Hodrick, R. J., Xing, Y., & Zhang, X. (2006). The cross-section of volatility and expected returns. *Journal of Finance*, 61(1), 259–299.

Banz, R. W. (1981). The relationship between return and market value of common stocks. *Journal of Financial Economics*, 9(1), 3–18.

Barbee, W. C., Mukherji, S., & Raines, G. A. (1996). Do sales-price and debt-equity explain stock returns better than book-market and firm size? *Financial Analysts Journal*, 52(2), 56–60.

Bhandari, L. C. (1988). Debt/equity ratio and expected common stock returns: Empirical evidence. *Journal of Finance*, 43(2), 507–528.

Cooper, M. J., Gulen, H., & Schill, M. J. (2008). Asset growth and the cross-section of stock returns. *Journal of Finance*, 63(4), 1609–1651.

Datar, V. T., Naik, N. Y., & Radcliffe, R. (1998). Liquidity and stock returns: An alternative test. *Journal of Financial Markets*, 1(2), 203–219.

Davis, J. L., Fama, E. F., & French, K. R. (2000). Characteristics, covariances, and average returns: 1929 to 1997. *Journal of Finance*, 55(1), 389–406.

DeBondt, W. F. M., & Thaler, R. (1985). Does the stock market overreact? *Journal of Finance*, 40(3), 793–805.

Enache, L., & Srivastava, A. (2018). Should intangible investments be reported separately or commingled with operating expenses? New evidence. *Management Science*, 64(7), 3446–3468.

Chen, L., Pelger, M., & Zhu, J. (2024). Deep learning in asset pricing. *Management Science*, 70(2), 714–750.

Fama, E. F., & French, K. R. (1992). The cross-section of expected stock returns. *Journal of Finance*, 47(2), 427–465.

Fama, E. F., & French, K. R. (1993). Common risk factors in the returns on stocks and bonds. *Journal of Financial Economics*, 33(1), 3–56.

Fama, E. F., & French, K. R. (2015). A five-factor asset pricing model. *Journal of Financial Economics*, 116(1), 1–22.

Freyberger, J., Neuhierl, A., & Weber, M. (2020). Dissecting characteristics nonparametrically. *Review of Financial Studies*, 33(5), 2326–2377.

George, T. J., & Hwang, C.-Y. (2004). The 52-week high and momentum investing. *Journal of Finance*, 59(5), 2145–2176.

Grinblatt, M., & Moskowitz, T. J. (2004). Predicting stock price movements from past returns: The role of consistency and tax-loss selling. *Journal of Financial Economics*, 71(3), 541–579.

Hirshleifer, D., Hou, K., Teoh, S. H., & Zhang, Y. (2004). Do investors overvalue firms with bloated balance sheets? *Journal of Accounting and Economics*, 38, 297–331.

Hou, K., Xue, C., & Zhang, L. (2015). Digesting anomalies: An investment approach. *Review of Financial Studies*, 28(3), 650–705.

Jegadeesh, N. (1990). Evidence of predictable behavior of security returns. *Journal of Finance*, 45(3), 881–898.

Jegadeesh, N., & Titman, S. (1993). Returns to buying winners and selling losers: Implications for stock market efficiency. *Journal of Finance*, 48(1), 65–91.

Lakonishok, J., Shleifer, A., & Vishny, R. W. (1994). Contrarian investment, extrapolation, and risk. *Journal of Finance*, 49(5), 1541–1578.

Llorente, G., Michaely, R., Saar, G., & Wang, J. (2002). Dynamic volume-return relation of individual stocks. *Review of Financial Studies*, 15(4), 1005–1047.

Lyandres, E., Sun, L., & Zhang, L. (2008). The new issues puzzle: Testing the investment-based explanation. *Review of Financial Studies*, 21(6), 2825–2855.

Novy-Marx, R. (2011). Operating leverage. *Review of Finance*, 15(1), 103–134.

Novy-Marx, R. (2013). The other side of value: The gross profitability premium. *Journal of Financial Economics*, 108(1), 1–28.

Palazzo, B. (2012). Cash holdings, risk, and expected returns. *Journal of Financial Economics*, 104(1), 162–185.

Pontiff, J., & Woodgate, A. (2008). Share issuance and cross-sectional returns. *Journal of Finance*, 63(2), 921–945.

Richardson, S. A., Sloan, R. G., Soliman, M. T., & Tuna, I. (2005). Accrual reliability, earnings persistence and stock prices. *Journal of Accounting and Economics*, 39(3), 437–485.

Roll, R. (1984). A simple implicit measure of the effective bid-ask spread in an efficient market. *Journal of Finance*, 39(4), 1127–1139.

Sharpe, W. F. (1964). Capital asset prices: A theory of market equilibrium under conditions of risk. *Journal of Finance*, 19(3), 425–442.

Shumway, T. (1997). The delisting bias in CRSP data. *Journal of Finance*, 52(1), 327–340.

Sloan, R. G. (1996). Do stock prices fully reflect information in accruals and cash flows about future earnings? *Accounting Review*, 71(3), 289–315.

Soliman, M. T. (2008). The use of DuPont analysis by market participants. *Accounting Review*, 83(3), 823–853.

Titman, S., Wei, K. C. J., & Xie, F. (2004). Capital investments and stock returns. *Journal of Financial and Quantitative Analysis*, 39(4), 677–700.